

# Final Project Submission

Please fill out:

- Student name: Marife Ramoran
  - Student pace: self paced
  - Scheduled project review date/time: 01 December 2023
  - Instructor name: Hardik Idhani
  - Blog post URL:
- 

## Overview

---

The undertaking serves as a valuable showcase for applying the knowledge acquired during Phase 2 to real-world data. By employing data analysis and linear regression, our objective is to derive a minimum of three key insights that can inform decision-making regarding the profitability of renovating older homes. A crucial aspect involves comprehending the factors influencing house prices.

## Business Problem

---

The real estate agency seeks to offer strategic advice to potential homebuyers planning to acquire and renovate properties. This advice aims to pinpoint essential renovation priorities that can maximize the property's market value growth, along with providing estimates of value increments to optimize their resale potential.

## Data

---

The dataset utilized for this model comprises information from houses in King County, which were sold between 2014 and 2015. It encompasses 21,597 entries, incorporating details about the house such as the number of bedrooms, bathrooms, and its rating (grade).

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
from statsmodels.formula.api import ols
```

```
In [2]: data = pd.read_csv("data/kc_house_data.csv")
```

```
In [3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21597 non-null  int64
1   date                  21597 non-null  object
2   price                 21597 non-null  float64
3   bedrooms              21597 non-null  int64
4   bathrooms             21597 non-null  float64
5   sqft_living           21597 non-null  int64
6   sqft_lot              21597 non-null  int64
7   floors                21597 non-null  float64
8   waterfront            19221 non-null  float64
9   view                  21534 non-null  float64
10  condition             21597 non-null  int64
11  grade                 21597 non-null  int64
12  sqft_above            21597 non-null  int64
13  sqft_basement         21597 non-null  object
14  yr_built              21597 non-null  int64
15  yr_renovated          17755 non-null  float64
16  zipcode               21597 non-null  int64
17  lat                   21597 non-null  float64
18  long                  21597 non-null  float64
19  sqft_living15         21597 non-null  int64
20  sqft_lot15            21597 non-null  int64
dtypes: float64(8), int64(11), object(2)
memory usage: 3.5+ MB
```

In [4]: `data.describe()`

Out [4]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lo
<b>count</b>	2.159700e+04	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04
<b>mean</b>	4.580474e+09	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04
<b>std</b>	2.876736e+09	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04
<b>min</b>	1.000102e+06	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+03
<b>25%</b>	2.123049e+09	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03
<b>50%</b>	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03
<b>75%</b>	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04
<b>max</b>	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+05

In [5]: `data.head()`

Out [5]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	wa
<b>0</b>	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	
<b>1</b>	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
<b>2</b>	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
<b>3</b>	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
<b>4</b>	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	

5 rows × 10 columns

```
In [6]: # I want to check the duplicates
data[data.duplicated(keep=False, subset=['id'])]
```

Out [6]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floor
<b>93</b>	6021501535	7/25/2014	430000.0	3	1.50	1580	5000	1.
<b>94</b>	6021501535	12/23/2014	700000.0	3	1.50	1580	5000	1.
<b>313</b>	4139480200	6/18/2014	1380000.0	4	3.25	4290	12103	1.
<b>314</b>	4139480200	12/9/2014	1400000.0	4	3.25	4290	12103	1.
<b>324</b>	7520000520	9/5/2014	232000.0	2	1.00	1240	12092	1.
...	...	...	...	...	...	...	...	.
<b>20654</b>	8564860270	3/30/2015	502000.0	4	2.50	2680	5539	2.
<b>20763</b>	6300000226	6/26/2014	240000.0	4	1.00	1200	2171	1.
<b>20764</b>	6300000226	5/4/2015	380000.0	4	1.00	1200	2171	1.
<b>21564</b>	7853420110	10/3/2014	594866.0	3	3.00	2780	6000	2.
<b>21565</b>	7853420110	5/4/2015	625000.0	3	3.00	2780	6000	2.

353 rows × 21 columns

```
In [7]: # Now, I am eliminating duplicates while preserving the initial row
data = data.drop_duplicates(keep= 'first', subset= ['id'])
```

To enhance data readability, we can omit the following variables, as they are not directly linked to the target variable "price":

- view: Indicates whether the house has been viewed.
- lat and long: Represent the coordinates of the house.
- zipcode: Unlikely to influence the outcome, as house prices post-renovation are not correlated with their zip codes.
- date: Indicates the sale date of the house.
- sqft\_lot15 and sqft\_living15: Reflect the size of other houses in the vicinity.
- id: Serves as the unique identification number for the listed house.
- yr\_built: indicates which year the house was built

```
In [8]: to_drop = ['view', 'lat', 'long', 'zipcode', 'date', 'sqft_lot15',
data = data.drop(to_drop, axis= 1)
```

```
In [9]: # Let's examine whether there are any null values present in our data
data.isna().sum()
```

```
Out[9]: price                0
bedrooms                   0
bathrooms                  0
sqft_living                0
sqft_lot                   0
floors                     0
waterfront                2353
condition                  0
grade                      0
sqft_above                 0
sqft_basement              0
yr_renovated               3804
dtype: int64
```

```
In [10]: # We discovered that the variable "sqft_basement" contains some unusual values
data['sqft_basement'].value_counts()
```

```
Out[10]: 0.0          12717
?          452
600.0       216
500.0       206
700.0       205
...
1920.0        1
3480.0        1
2730.0        1
2720.0        1
248.0         1
Name: sqft_basement, Length: 304, dtype: int64
```

```
In [11]: # We identified 454 instances of missing values in the "sqft_basement" variable
# To address this, we will replace the "?" entries with 0 to render the data in a consistent format
data['sqft_basement'] = data['sqft_basement'].replace('?',0)
```

```
In [12]: # Observing that the "waterfront" variable contains some missing data
data['waterfront'].unique()
```

```
Out[12]: array([nan,  0.,  1.])
```

```
In [13]: # Upon examination, we have identified missing data that can be filled in
data['waterfront'].fillna(value = 0, inplace = True)
```

```
In [14]: # We've also identified missing data in the "yr_renovated" variable
data['yr_renovated'].unique()
```

```
Out[14]: array([ 0., 1991., nan, 2002., 2010., 1992., 2013., 1994., 1978.,
        2005., 2003., 1984., 1954., 2014., 2011., 1983., 1945., 1990.,
        1988., 1977., 1981., 1995., 2000., 1999., 1998., 1970., 1989.,
        2004., 1986., 2007., 1987., 2006., 1985., 2001., 1980., 1971.,
        1979., 1997., 1950., 1969., 1948., 2009., 2015., 1974., 2008.,
        1968., 2012., 1963., 1951., 1962., 1953., 1993., 1996., 1955.,
        1982., 1956., 1940., 1976., 1946., 1975., 1964., 1973., 1957.,
        1959., 1960., 1967., 1965., 1934., 1972., 1944., 1958.] )
```

```
In [15]: # Upon examination, we have detected missing data that can be filled
data['yr_renovated'].fillna(value = 0, inplace = True)
```

```
In [16]: # Upon closer examination, using float in models is not optimal.
# As a solution, we will multiply the bathroom values by 100, allowing
data['bathrooms'] = data['bathrooms'] * 100
```

```
In [17]: # We must convert the data into either integers or floats to facilitate
data['waterfront'] = data['waterfront'].astype(int)
data['bathrooms'] = data['bathrooms'].astype(int)
data['yr_renovated'] = data['yr_renovated'].astype(float).astype(int)
data['floors'] = data['floors'].astype(float).astype(int)
data['sqft_basement'] = data['sqft_basement'].astype(float).astype(int)
```

```
In [18]: data.describe()
```

```
Out[18]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
count	2.142000e+04	21420.000000	21420.000000	21420.000000	2.142000e+04	21420.000000
mean	5.407393e+05	3.373950	211.842904	2083.132633	1.512804e+04	1.447995
std	3.679311e+05	0.925405	76.871996	918.808412	4.153080e+04	0.552110
min	7.800000e+04	1.000000	50.000000	370.000000	5.200000e+02	1.000000
25%	3.225000e+05	3.000000	175.000000	1430.000000	5.040000e+03	1.000000
50%	4.500000e+05	3.000000	225.000000	1920.000000	7.614000e+03	1.000000
75%	6.450000e+05	4.000000	250.000000	2550.000000	1.069050e+04	2.000000
max	7.700000e+06	33.000000	800.000000	13540.000000	1.651359e+06	3.000000

```
In [19]: # We aim to streamline the data by transforming the information about the year
# Instead of specifying the year, we will simplify it to a straight yes/no
check = lambda x : x != 0
data['reno_bool'] = check(data['yr_renovated'])
```

```
In [20]: # Since many houses lack a basement, we can modify the variable to a flag
data['base_flag'] = check (data['sqft_basement'])
```

```
In [21]: data['reno_bool'] = data['reno_bool'].astype(int)
data['base_flag'] = data['base_flag'].astype(int)
```

```
In [22]: # Now, we can eliminate the columns used in creating our new variables
data = data.drop(['sqft_basement', 'yr_renovated'], axis=1)
```

```
In [23]: data.head()
```

Out[23]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	condition	grade
0	221900.0	3	100	1180	5650	1	0	3	7
1	538000.0	3	225	2570	7242	2	0	3	7
2	180000.0	2	100	770	10000	1	0	3	6
3	604000.0	4	300	1960	5000	1	0	5	7
4	510000.0	3	200	1680	8080	1	0	3	8

## Examine the outliers

We must identify and exclude any outliers in our dataset to enhance its applicability.

In [24]: *# Examine the distribution of the base through percentiles and medi*

```
fig= plt.figure(figsize=(18,10))

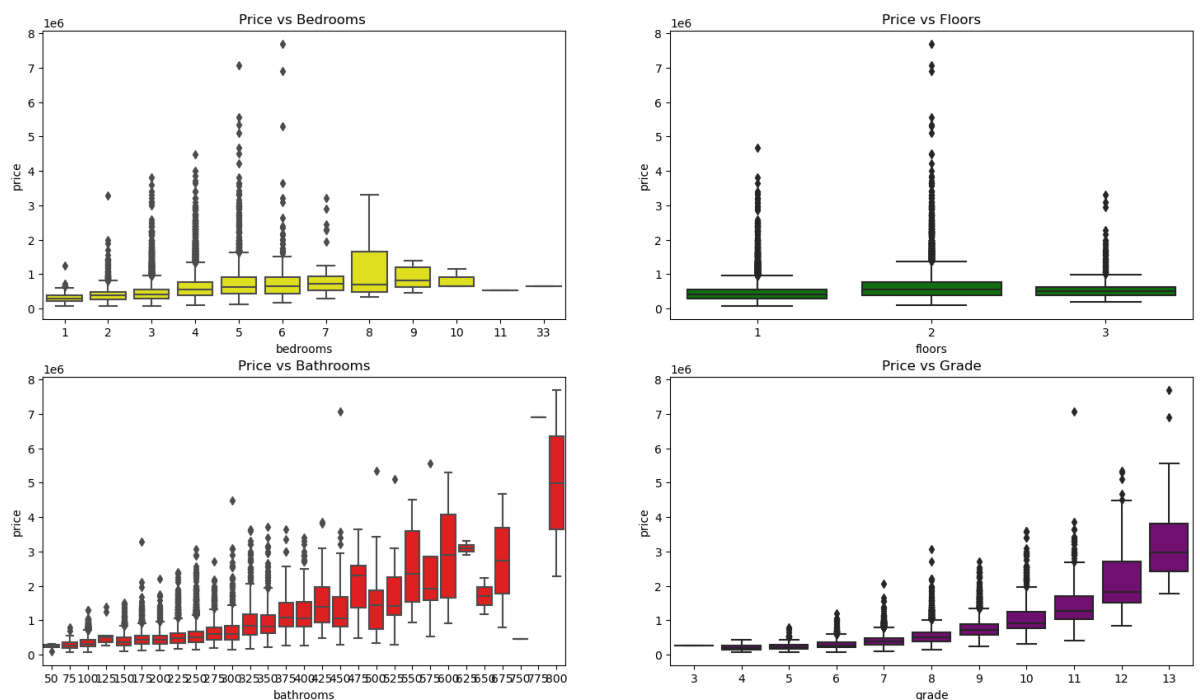
ax=fig.add_subplot(2,2,1)
sns.boxplot(data=data, x=data["bedrooms"], y=data["price"], hue=None)
ax.set_title("Price vs Bedrooms ")

ax=fig.add_subplot(2,2,2)
sns.boxplot(data=data, x=data["floors"], y=data["price"], hue=None,
ax.set_title("Price vs Floors")

ax=fig.add_subplot(2,2,3)
sns.boxplot(data=data, x=data["bathrooms"], y=data["price"], hue=None)
ax.set_title("Price vs Bathrooms")

ax=fig.add_subplot(2,2,4)
sns.boxplot(data=data, x=data["grade"], y=data["price"], hue=None,
ax.set_title("Price vs Grade")

plt.savefig(".\images\outliers.png", dpi = 150, bbox_inches = 'tight')
plt.show()
```





```
In [25]: # We will include all variables except for "reno_bool" and "waterfr
# This is essential to assess whether renovations can influence pri

columns= ['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lo
outliers=[]

for col in columns:
    mean=np.mean(data[col])
    std=np.std(data[col])
    for i in range(len(data)):
        item = data[col].iloc[i]
        z_score=(item-mean)/std
        if np.abs(z_score)>3:
            outliers.append(i)

outliers = set(outliers)
outliers = list(outliers)

# print (col)
```

```
In [26]: # We can now eliminate the outliers to create a cleaner dataset.
data.drop(data.index[outliers], inplace= True)
```

```

In [27]: # I want to check the boxplot once more for the outliers.

# Examine the distribution of the base through percentiles and medi

fig= plt.figure(figsize=(18,10))

ax=fig.add_subplot(2,2,1)
sns.boxplot(data=data, x=data["bedrooms"], y=data["price"], hue=None,
ax.set_title("Price vs bedrooms ")

ax=fig.add_subplot(2,2,2)
sns.boxplot(data=data, x=data["floors"], y=data["price"], hue=None,
ax.set_title("Price vs floors")

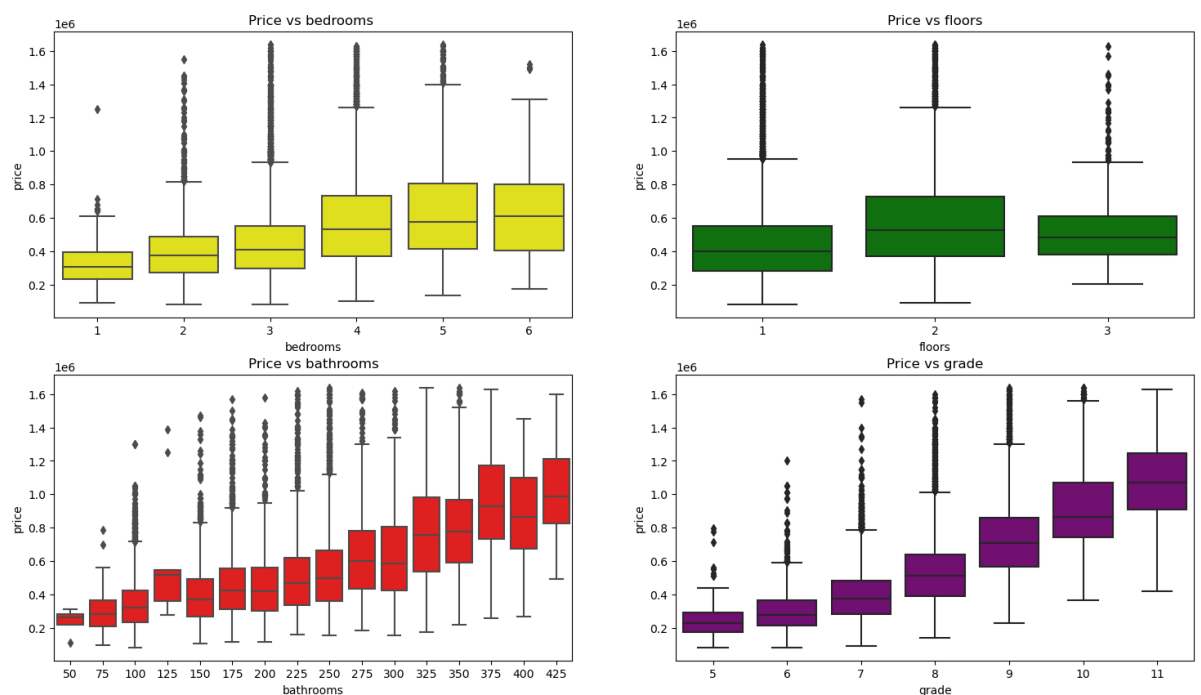
ax=fig.add_subplot(2,2,3)
sns.boxplot(data=data, x=data["bathrooms"], y=data["price"], hue=None,
ax.set_title("Price vs bathrooms")

ax=fig.add_subplot(2,2,4)
sns.boxplot(data=data, x=data["grade"], y=data["price"], hue=None,
ax.set_title("Price vs grade")

plt.savefig(".\images\outliers.png", dpi = 150, bbox_inches = 'tight')

plt.show()

```



# Model 1

```
In [28]: outcome = 'price'
predictors = data.drop(['price'], axis=1)
pred_sum = "+".join(predictors.columns)
formula = outcome + "~" + pred_sum
```

```
In [29]: model = ols(formula=formula, data=data).fit()
model.summary()
```

Out [29]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.533
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.532
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	2105.
<b>Date:</b>	Fri, 01 Dec 2023	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	16:25:58	<b>Log-Likelihood:</b>	-2.7391e+05
<b>No. Observations:</b>	20341	<b>AIC:</b>	5.478e+05
<b>Df Residuals:</b>	20329	<b>BIC:</b>	5.479e+05
<b>Df Model:</b>	11		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-6.495e+05	1.39e+04	-46.671	0.000	-6.77e+05	-6.22e+05
<b>bedrooms</b>	-2.046e+04	1825.138	-11.211	0.000	-2.4e+04	-1.69e+04
<b>bathrooms</b>	-211.2116	28.452	-7.423	0.000	-266.980	-155.444
<b>sqft_living</b>	127.1091	5.613	22.646	0.000	116.107	138.111
<b>sqft_lot</b>	-1.0910	0.094	-11.574	0.000	-1.276	-0.906
<b>floors</b>	-5711.7855	3113.702	-1.834	0.067	-1.18e+04	391.321
<b>waterfront</b>	3.376e+05	2.01e+04	16.784	0.000	2.98e+05	3.77e+05
<b>condition</b>	5.304e+04	1979.682	26.792	0.000	4.92e+04	5.69e+04
<b>grade</b>	1.077e+05	1808.823	59.537	0.000	1.04e+05	1.11e+05
<b>sqft_above</b>	3.4359	6.058	0.567	0.571	-8.439	15.310
<b>reno_bool</b>	1.315e+05	6807.655	19.315	0.000	1.18e+05	1.45e+05
<b>base_flag</b>	4.685e+04	4520.509	10.364	0.000	3.8e+04	5.57e+04

Omnibus = 2218.807 Prob (Chi-Square) = 1.000

<b>Omnibus:</b>	3516.397	<b>Durbin-Watson:</b>	1.960
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	7929.343
<b>Skew:</b>	0.930	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	5.429	<b>Cond. No.</b>	2.90e+05

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.9e+05. This might indicate that there are strong multicollinearity or other numerical problems.

## Summary on Model 1

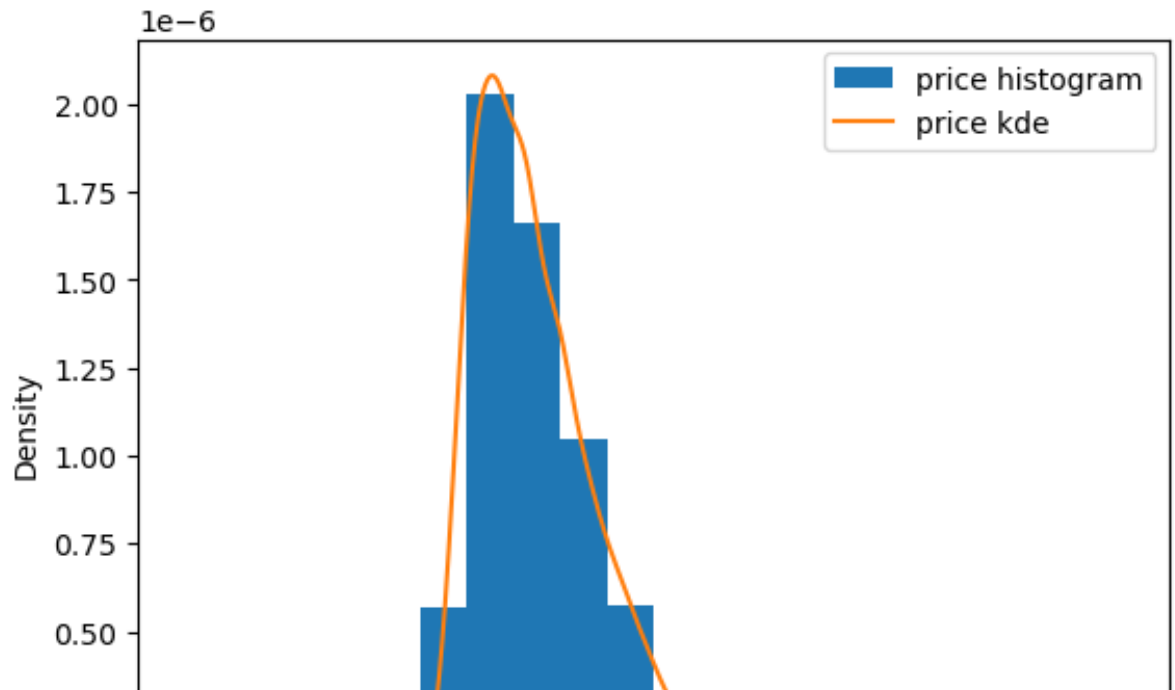
Currently, the P-values suggest that all independent variables, except for `sqft_above`, are statistically significant. Although the adjusted R-squared is satisfactory, there is room for improvement through the handling of categorical data, log transformation, and normalization.

The skewness at 0.930 indicates a positive skew, while the kurtosis at 5.429 signifies a leptokurtic curve. It suggests the presence of outliers, but these may be necessary for extracting valuable information from certain columns and rendering them usable.

## Checking assumptions

---

```
In [30]: # KDE
for column in data:
    data[column].plot.hist(density=True, label = column + ' histogram')
    data[column].plot.kde(label = column + ' kde')
plt.legend()
plt.show()
```



## Dealing with Categorical Data

---

```
In [31]: # We should create a histogram to confirm whether the data at hand

fig, axs = plt.subplots(3, 3, figsize=(12, 12))

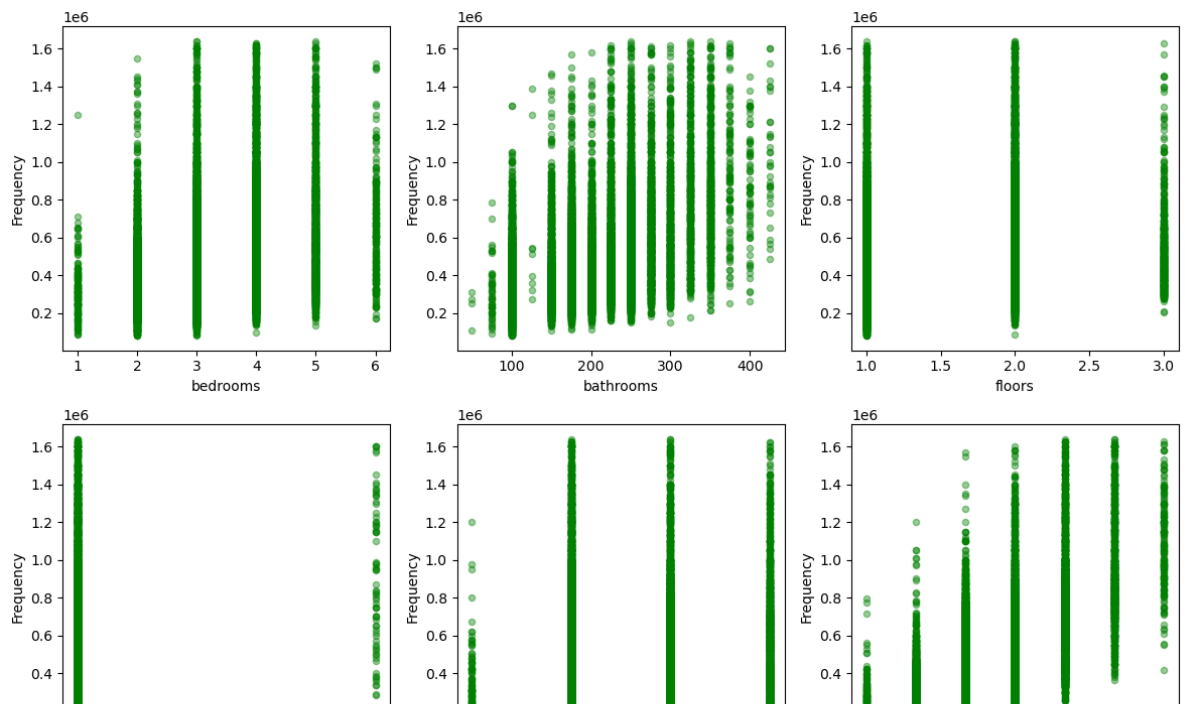
columns = ['bedrooms', 'bathrooms', 'floors', 'waterfront', 'condit

for i, column in enumerate(columns):
    row = i // 3
    col = i % 3
    data.plot(kind='scatter', x=column, y='price', ax=axs[row, col])
    axs[row, col].set_xlabel(column)
    axs[row, col].set_ylabel('Frequency')

plt.tight_layout()

plt.savefig(".\images\categorical.png", dpi = 150, bbox_inches = 't

plt.show()
```



Observing the clustered distribution in the scatter plot for the columns:

- bedrooms
- bathrooms
- floors
- waterfront
- condition
- grade
- reno\_bool
- base\_flag

It becomes evident that these are categorical data. To utilize them effectively, we should create dummy variables.

```
In [32]: categorical = ['bedrooms', 'bathrooms', 'floors', 'waterfront', 'co
```

```
In [33]: bed_dummies = pd.get_dummies(data['bedrooms'], prefix='bed', drop_f
bath_dummies = pd.get_dummies(data['bathrooms'], prefix='bath', drop
flr_dummies = pd.get_dummies(data['floors'], prefix='flr', drop_fir
wtr_dummies = pd.get_dummies(data['waterfront'], prefix='wtr', drop
cond_dummies = pd.get_dummies(data['condition'], prefix='cond', drop
grade_dummies = pd.get_dummies(data['grade'], prefix='grade', drop_
reno_dummies = pd.get_dummies(data['reno_bool'], prefix='reno', drop
base_dummies = pd.get_dummies(data['base_flag'], prefix='base', drop
```

```
In [34]: dummies = pd.concat([bed_dummies, bath_dummies, flr_dummies, wtr_du
```

```
In [35]: bed_dummies = pd.get_dummies(data['bedrooms'], prefix='bed')
bath_dummies = pd.get_dummies(data['bathrooms'], prefix='bath')
flr_dummies = pd.get_dummies(data['floors'], prefix='flr')
wtr_dummies = pd.get_dummies(data['waterfront'], prefix='wtr')
cond_dummies = pd.get_dummies(data['condition'], prefix='cond')
grade_dummies = pd.get_dummies(data['grade'], prefix='grade')
reno_dummies = pd.get_dummies(data['reno_bool'], prefix='reno')

base_dummies = pd.get_dummies(data['base_flag'], prefix='base')
```

```
In [36]: dummies_full = pd.concat([bed_dummies, bath_dummies, flr_dummies, w
```

```
In [37]: data_1 = data.drop(categorical, axis=1)
```

## Multicollinearity

To enhance model performance and ensure accurate coefficients, it is essential to eliminate highly correlated variables.

```
In [38]: data_pred = data_1.iloc[:, 1:]
```


```
In [39]: sm = pd.plotting.scatter_matrix(data_pred, figsize = [16, 16]);

# Rotate the text
[s.xaxis.label.set_rotation(90) for s in sm.reshape(-1)]
[s.yaxis.label.set_rotation(0) for s in sm.reshape(-1)]

# Consider adjusting the label offset when rotating to avoid overlap
[s.get_yaxis().set_label_coords(-1, 0.5) for s in sm.reshape(-1)]

# Conceal the ticks
[s.set_xticks(()) for s in sm.reshape(-1)]
[s.set_yticks(()) for s in sm.reshape(-1)]

plt.show()
```





```
In [40]: # Present the correlation in numerical format.
data_pred.corr()
```

Out[40]:

	sqft_living	sqft_lot	sqft_above
sqft_living	1.000000	0.219082	0.84521
sqft_lot	0.219082	1.000000	0.20770
sqft_above	0.845210	0.207700	1.00000

```
In [41]: abs(data_pred.corr()) > 0.75
```

Out[41]:

	sqft_living	sqft_lot	sqft_above
sqft_living	True	False	True
sqft_lot	False	True	False
sqft_above	True	False	True

```
In [42]: # This will display pairs with identical correlation values, prevent
df=data_pred.corr().abs().stack().reset_index().sort_values(0, ascending=False)

# Create a new column named "pairs" by zipping the variable name columns
# which were initially named "level_0" and "level_1" by default.
df['pairs'] = list(zip(df.level_0, df.level_1))

# Set the index to pairs
df.set_index(['pairs'], inplace = True)

# Drop the columns in "levels"
df.drop(columns=['level_1', 'level_0'], inplace = True)

# Rename the correlation column to "cc" instead of "0".
df.columns = ['cc']

# Remove duplicates.
# Exercise caution, especially if you have variables that are perfectly correlated
# as this action could be risky.
df.drop_duplicates(inplace=True)
```

```
In [43]: df[(df.cc>0.75) & (df.cc < 1.0)]
```

Out[43]:

	cc
pairs	
(sqft_living, sqft_above)	0.84521

## Dealing with Multicollinearity

We identified a high correlation between the variables "sqft\_above" and "sqft\_living." It is appropriate to address this by dropping one variable from the pair. In this case, we will eliminate "sqft\_above," given that its description corresponds to the square footage of the house.

```
In [44]: dt = data_1.drop(['sqft_above'], axis= 1)
```

```
In [45]: dt.head()
```

Out [45]:

	price	sqft_living	sqft_lot
0	221900.0	1180	5650
1	538000.0	2570	7242
2	180000.0	770	10000
3	604000.0	1960	5000
4	510000.0	1680	8080

## Model 2

```
In [46]: dt_1_dummies = pd.concat([dt, dummies], axis= 1)
```

```
In [47]: dt_1_dummies.head()
```

Out [47]:

	price	sqft_living	sqft_lot	bed_2	bed_3	bed_4	bed_5	bed_6	bath_75	bath_100	.
0	221900.0	1180	5650	0	1	0	0	0	0	1	.
1	538000.0	2570	7242	0	1	0	0	0	0	0	.
2	180000.0	770	10000	1	0	0	0	0	0	1	.
3	604000.0	1960	5000	0	0	1	0	0	0	0	.
4	510000.0	1680	8080	0	1	0	0	0	0	0	.

5 rows × 37 columns

```
In [48]: predictors = dt_1_dummies.drop(['price'], axis = 1)
pred_sum = "+".join(predictors.columns)
formula = outcome + "~" + pred_sum
```

```
In [49]: model = ols(formula = formula, data = dt_1_dummies).fit()
model.summary()
```

Out [49]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.557			
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.556			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	707.7			
<b>Date:</b>	Fri, 01 Dec 2023	<b>Prob (F-statistic):</b>	0.00			
<b>Time:</b>	16:26:03	<b>Log-Likelihood:</b>	-2.7337e+05			
<b>No. Observations:</b>	20341	<b>AIC:</b>	5.468e+05			
<b>Df Residuals:</b>	20304	<b>BIC:</b>	5.471e+05			
<b>Df Model:</b>	36					
<b>Covariance Type:</b>	nonrobust					
	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>[0.025</b>	<b>0.975]</b>
<b>Intercept</b>	-8880.3248	8.52e+04	-0.104	0.917	-1.76e+05	1.58e+05
<b>sqft_living</b>	120.6559	3.071	39.284	0.000	114.636	126.676
<b>sqft_lot</b>	-1.0623	0.092	-11.488	0.000	-1.244	-0.881
<b>bed_2</b>	-6399.8661	1.35e+04	-0.474	0.635	-3.29e+04	2.01e+04
<b>bed_3</b>	-5.026e+04	1.35e+04	-3.729	0.000	-7.67e+04	-2.38e+04
<b>bed_4</b>	-5.679e+04	1.38e+04	-4.127	0.000	-8.38e+04	-2.98e+04
<b>bed_5</b>	-5.805e+04	1.45e+04	-3.999	0.000	-8.65e+04	-2.96e+04
<b>bed_6</b>	-6.037e+04	1.82e+04	-3.313	0.001	-9.61e+04	-2.47e+04
<b>bath_75</b>	9.278e+04	8.61e+04	1.077	0.281	-7.61e+04	2.62e+05
<b>bath_100</b>	1.168e+05	8.33e+04	1.403	0.161	-4.64e+04	2.8e+05
<b>bath_125</b>	1.383e+05	1e+05	1.382	0.167	-5.78e+04	3.34e+05
<b>bath_150</b>	9.195e+04	8.34e+04	1.103	0.270	-7.15e+04	2.55e+05
<b>bath_175</b>	8.647e+04	8.33e+04	1.038	0.299	-7.69e+04	2.5e+05
<b>bath_200</b>	9.046e+04	8.34e+04	1.085	0.278	-7.3e+04	2.54e+05
<b>bath_225</b>	7.772e+04	8.34e+04	0.932	0.351	-8.57e+04	2.41e+05
<b>bath_250</b>	4.87e+04	8.34e+04	0.584	0.559	-1.15e+05	2.12e+05
<b>bath_275</b>	7.819e+04	8.35e+04	0.937	0.349	-8.54e+04	2.42e+05

<b>bath_300</b>	8.902e+04	8.36e+04	1.065	0.287	-7.48e+04	2.53e+05
<b>bath_325</b>	1.167e+05	8.37e+04	1.393	0.164	-4.74e+04	2.81e+05
<b>bath_350</b>	1.043e+05	8.37e+04	1.247	0.212	-5.97e+04	2.68e+05
<b>bath_375</b>	1.943e+05	8.5e+04	2.286	0.022	2.77e+04	3.61e+05
<b>bath_400</b>	1.016e+05	8.58e+04	1.184	0.236	-6.66e+04	2.7e+05
<b>bath_425</b>	1.778e+05	8.84e+04	2.012	0.044	4601.006	3.51e+05
<b>flr_2</b>	-1.191e+04	3633.602	-3.277	0.001	-1.9e+04	-4785.042
<b>flr_3</b>	5.127e+04	7793.927	6.578	0.000	3.6e+04	6.65e+04
<b>wtr_1</b>	3.257e+05	1.97e+04	16.547	0.000	2.87e+05	3.64e+05
<b>cond_3</b>	1.631e+04	1.39e+04	1.176	0.240	-1.09e+04	4.35e+04
<b>cond_4</b>	6.337e+04	1.39e+04	4.548	0.000	3.61e+04	9.07e+04
<b>cond_5</b>	1.302e+05	1.44e+04	9.061	0.000	1.02e+05	1.58e+05
<b>grade_6</b>	3.825e+04	1.2e+04	3.187	0.001	1.47e+04	6.18e+04
<b>grade_7</b>	1.118e+05	1.18e+04	9.433	0.000	8.85e+04	1.35e+05
<b>grade_8</b>	2.153e+05	1.22e+04	17.612	0.000	1.91e+05	2.39e+05
<b>grade_9</b>	3.648e+05	1.29e+04	28.366	0.000	3.4e+05	3.9e+05
<b>grade_10</b>	4.841e+05	1.4e+04	34.517	0.000	4.57e+05	5.12e+05
<b>grade_11</b>	5.988e+05	1.84e+04	32.607	0.000	5.63e+05	6.35e+05
<b>reno_1</b>	1.251e+05	6666.666	18.761	0.000	1.12e+05	1.38e+05
<b>base_1</b>	4.481e+04	2834.264	15.810	0.000	3.93e+04	5.04e+04

<b>Omnibus:</b>	3249.639	<b>Durbin-Watson:</b>	1.964
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	8265.134
<b>Skew:</b>	0.892	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	5.564	<b>Cond. No.</b>	4.91e+06

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 4.91e+06. This might indicate that there are strong multicollinearity or other numerical problems.

## Summary on Model 2

Addressing multicollinearity resulted in an increase in the R-squared value, indicating that we are moving in the right direction.

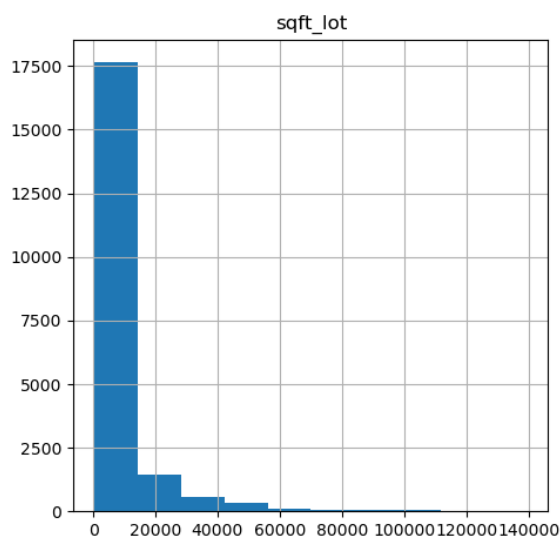
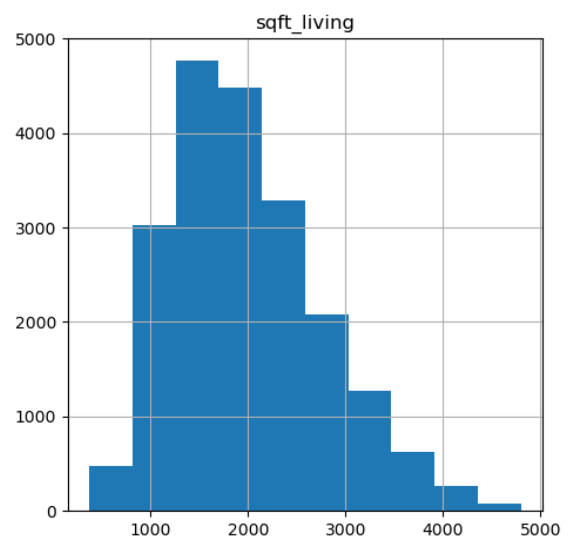

## Applying log transformation and normalization

In [50]: *# We aim to visualize our data through a histogram.  
# Upon examination, I propose performing a log transformation and n  
# to bring the mean, median, and mode closer to zero.*

```
fig = plt.figure(figsize = (12,12))
ax = fig.gca()
dt.hist(ax = ax);
```


/var/folders/zr/m5b0ldvs7695p\_h8w70ndwc40000gn/T/ipykernel\_57858/3211950891.py:7: UserWarning: To output multiple subplots, the figure containing the passed axes is being cleared.

```
dt.hist(ax = ax);
```



In [51]: *# Let's examine the data after the transformation.*

```
data_log = pd.DataFrame([])
data_log['sqft_living'] = np.log(dt['sqft_living'])
data_log['sqft_lot'] = np.log(dt['sqft_lot'])
data_log.hist(figsize = [6, 6]);
```



```
In [52]: log_living = data_log['sqft_living']
log_lot = data_log['sqft_lot']

# We can normalize our data


scaled_living = (log_living - np.mean(log_living)) / np.sqrt(np.var(log_living))
scaled_lot = (log_lot - np.mean(log_lot)) / np.sqrt(np.var(log_lot))
```

In [53]: *# Now, let's assess the current state of our data.*

```
data_cont_scaled = pd.DataFrame([])

data_cont_scaled['sqft_living'] = scaled_living
data_cont_scaled['sqft_lot'] = scaled_lot

data_cont_scaled.hist(figsize = [6, 6]);
```



## Model 3

In [54]: *# Let's evaluate our model*

```
dt_1 = pd.concat([dt['price'], data_cont_scaled, dummies], axis = 1)
```

```
In [55]: predictors = dt_1.drop(['price'], axis = 1)
pred_sum = "+".join(predictors.columns)
formula = outcome + "~" + pred_sum
```

```
In [56]: model = ols(formula = formula, data = dt_1).fit()
model.summary()
```

Out [56]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.558			
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.557			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	712.6			
<b>Date:</b>	Fri, 01 Dec 2023	<b>Prob (F-statistic):</b>	0.00			
<b>Time:</b>	16:26:04	<b>Log-Likelihood:</b>	-2.7333e+05			
<b>No. Observations:</b>	20341	<b>AIC:</b>	5.467e+05			
<b>Df Residuals:</b>	20304	<b>BIC:</b>	5.470e+05			
<b>Df Model:</b>	36					
<b>Covariance Type:</b>	nonrobust					
	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>[0.025</b>	<b>0.975]</b>
<b>Intercept</b>	3.123e+05	8.52e+04	3.665	0.000	1.45e+05	4.79e+05
<b>sqft_living</b>	1.014e+05	2529.251	40.075	0.000	9.64e+04	1.06e+05
<b>sqft_lot</b>	-2.901e+04	1497.178	-19.374	0.000	-3.19e+04	-2.61e+04
<b>bed_2</b>	-3.283e+04	1.35e+04	-2.431	0.015	-5.93e+04	-6355.299
<b>bed_3</b>	-8.306e+04	1.36e+04	-6.121	0.000	-1.1e+05	-5.65e+04
<b>bed_4</b>	-8.775e+04	1.39e+04	-6.304	0.000	-1.15e+05	-6.05e+04
<b>bed_5</b>	-8.22e+04	1.47e+04	-5.603	0.000	-1.11e+05	-5.34e+04
<b>bed_6</b>	-8.219e+04	1.83e+04	-4.482	0.000	-1.18e+05	-4.62e+04
<b>bath_75</b>	1.098e+05	8.6e+04	1.276	0.202	-5.88e+04	2.78e+05
<b>bath_100</b>	1.174e+05	8.31e+04	1.412	0.158	-4.55e+04	2.8e+05
<b>bath_125</b>	1.049e+05	9.98e+04	1.051	0.293	-9.08e+04	3.01e+05
<b>bath_150</b>	7.751e+04	8.32e+04	0.931	0.352	-8.56e+04	2.41e+05
<b>bath_175</b>	6.968e+04	8.32e+04	0.838	0.402	-9.34e+04	2.33e+05
<b>bath_200</b>	6.938e+04	8.32e+04	0.834	0.404	-9.37e+04	2.32e+05
<b>bath_225</b>	6.295e+04	8.32e+04	0.756	0.449	-1e+05	2.26e+05
<b>bath_250</b>	3.177e+04	8.32e+04	0.382	0.703	-1.31e+05	1.95e+05
<b>bath_275</b>	6.539e+04	8.33e+04	0.785	0.433	-9.79e+04	2.29e+05



<b>bath_300</b>	7.558e+04	8.34e+04	0.906	0.365	-8.79e+04	2.39e+05
<b>bath_325</b>	1.125e+05	8.36e+04	1.346	0.178	-5.13e+04	2.76e+05
<b>bath_350</b>	1.054e+05	8.35e+04	1.262	0.207	-5.83e+04	2.69e+05
<b>bath_375</b>	1.987e+05	8.48e+04	2.343	0.019	3.24e+04	3.65e+05
<b>bath_400</b>	1.124e+05	8.56e+04	1.313	0.189	-5.54e+04	2.8e+05
<b>bath_425</b>	2.025e+05	8.82e+04	2.297	0.022	2.97e+04	3.75e+05
<b>flr_2</b>	-3.129e+04	3815.818	-8.201	0.000	-3.88e+04	-2.38e+04
<b>flr_3</b>	-4473.8620	8380.682	-0.534	0.593	-2.09e+04	1.2e+04
<b>wtr_1</b>	3.378e+05	1.97e+04	17.159	0.000	2.99e+05	3.76e+05
<b>cond_3</b>	8779.6520	1.38e+04	0.634	0.526	-1.84e+04	3.59e+04
<b>cond_4</b>	5.88e+04	1.39e+04	4.228	0.000	3.15e+04	8.61e+04
<b>cond_5</b>	1.225e+05	1.44e+04	8.539	0.000	9.44e+04	1.51e+05
<b>grade_6</b>	1.943e+04	1.2e+04	1.620	0.105	-4081.559	4.29e+04
<b>grade_7</b>	7.796e+04	1.19e+04	6.548	0.000	5.46e+04	1.01e+05
<b>grade_8</b>	1.813e+05	1.23e+04	14.702	0.000	1.57e+05	2.05e+05
<b>grade_9</b>	3.457e+05	1.29e+04	26.700	0.000	3.2e+05	3.71e+05
<b>grade_10</b>	4.871e+05	1.4e+04	34.769	0.000	4.6e+05	5.15e+05
<b>grade_11</b>	6.22e+05	1.82e+04	34.110	0.000	5.86e+05	6.58e+05
<b>reno_1</b>	1.254e+05	6654.320	18.842	0.000	1.12e+05	1.38e+05
<b>base_1</b>	2.955e+04	2960.604	9.980	0.000	2.37e+04	3.54e+04

<b>Omnibus:</b>	3280.745	<b>Durbin-Watson:</b>	1.963
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	8404.741
<b>Skew:</b>	0.897	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	5.588	<b>Cond. No.</b>	465.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

## Summary of Model 3 and Comparison

Following the transformation of our data and reintroduction of categorical variables, the model appears robust. The R-squared value of 0.558 is considered good, suggesting that the model is likely at its optimal fit. We can now proceed with model validation.

## Validation

---

```
In [57]: from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
```

```
In [58]: y = np.log(dt_1['price'])
x = dt_1.drop('price', axis = 1)
```

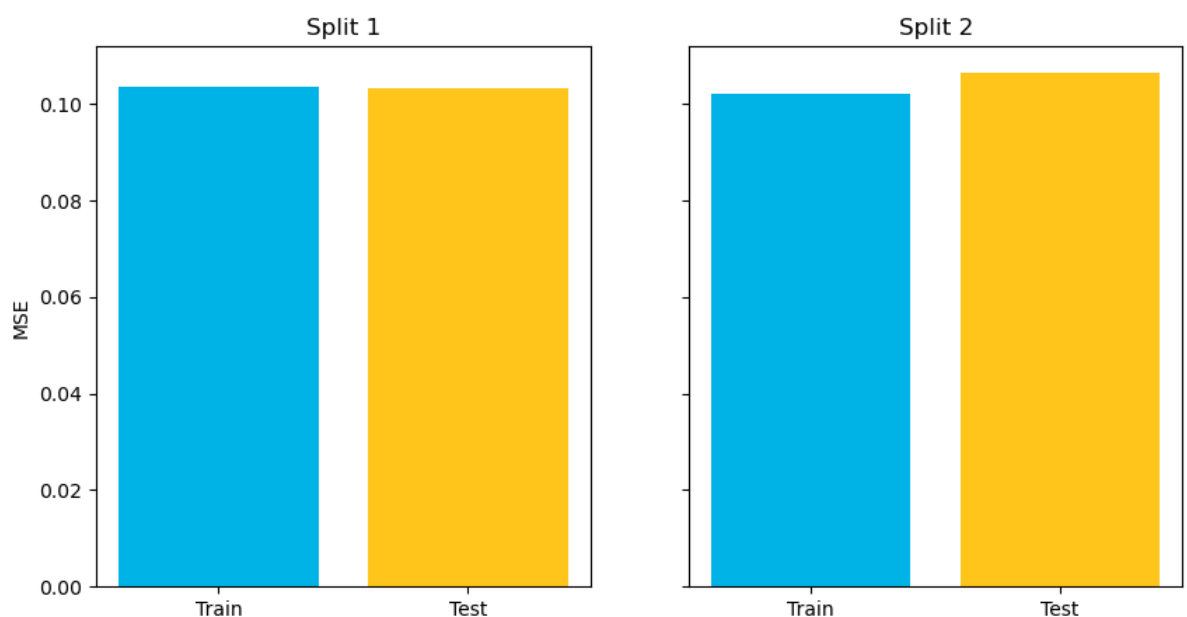
In [59]: *# Configure data and model set up*

```
linreg = LinearRegression()

# Split with random_state 87
x_train_1, x_test_1, y_train_1, y_test_1 = train_test_split(x, y, r
linreg.fit(x_train_1, y_train_1)
train_mse_1 = mean_squared_error(y_train_1, linreg.predict(x_train_
test_mse_1 = mean_squared_error(y_test_1, linreg.predict(x_test_1))

# Split with random_state 41
x_train_2, x_test_2, y_train_2, y_test_2 = train_test_split(x, y, r
linreg.fit(x_train_2, y_train_2)
train_mse_2 = mean_squared_error(y_train_2, linreg.predict(x_train_
test_mse_2 = mean_squared_error(y_test_2, linreg.predict(x_test_2))

# Plot metrics
fig, (left, right) = plt.subplots(ncols=2, figsize=(10,5), sharey=T
labels = ["Train", "Test"]
colors = ["#00B3E6", "#FFC51B"]
left.bar(labels, [train_mse_1, test_mse_1], color=colors)
left.set_title("Split 1")
left.set_ylabel("MSE")
right.bar(labels, [train_mse_2, test_mse_2], color=colors)
right.set_title("Split 2");
```



## Train-Test Split Analysis

It is evident that the train and test splits for both random\_state 87 and random\_state 41 exhibit only minimal differences.

```
In [60]: # Now, let's evaluate it with random_state 4
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=
```

```
In [61]: linreg = LinearRegression()
linreg.fit(x_train, y_train)
LinearRegression()
```

```
Out [61]: ▾ LinearRegression
LinearRegression()
```

## Calculate the Mean Squared Error

```
In [62]: linreg = LinearRegression()
linreg.fit(x_train, y_train)

y_hat_train = linreg.predict(x_train)
y_hat_test = linreg.predict(x_test)
```

```
In [63]: train_residuals = y_hat_train - y_train
test_residuals = y_hat_test - y_test
```

```
In [64]: train_mse = mean_squared_error(y_train, y_hat_train)
test_mse = mean_squared_error(y_test, y_hat_test)
print('Train Mean Squared Error:', train_mse)
print('Test Mean Squared Error:', test_mse)
```

Train Mean Squared Error: 0.10386598952456463  
Test Mean Squared Error: 0.10185845860770103

```
In [65]: train_mse - test_mse
```

```
Out [65]: 0.0020075309168636063
```

## R-squared value of the model

```
In [66]: print('R2 score: {:.2%}'.format(r2_score(y_test, y_hat_test)))

R2 score: 55.14%
```

The discrepancy is so minimal that rounding it to two decimal places results in 0.00. This suggests that the model is likely at its optimal fit.

```
from sklearn.feature_selection import RFE

selector = RFE(linreg, n_features_to_select=3)
selector = selector.fit(x_train, y_train.values)
selector.support_
```

```
selected_columns = x_train.columns[selector.support_]
linreg.fit(x_train[selected_columns], y_train)
print(selected_columns)
```

---

```
Index(['grade_9', 'grade_10', 'grade_11'], dtype='object')
```

## Evaluation

---

After three iterations, we discovered that the final model utilized the top independent variables. Using sklearn's feature selector, we identified variables with a strong relationship with the dependent variable, price. The top variables influencing the price were:

- grade\_11
- grade\_10
- grade\_9

Addressing the key renovation properties to maximize the property's market value, we identified the top five contributors:

1. sqft\_living: Increasing the square footage of the home can raise the price by 83,010.
2. flr\_3: Having a three-story house would increase the property's worth by 66,690.
3. bath\_425: Adding 4.25 bathrooms can increase the property's worth by 304,400.
4. wtr\_1: Having a waterfront view will increase the property's worth by 305,700.
5. grade\_11: Achieving a building grade of 11 in King County can increase the property's worth by 714,500.

This linear regression model effectively addresses the business problem of identifying key renovation properties to maximize the property's market value.

While the model provides meaningful insights, validating the outliers with more data could enhance its accuracy. It is noteworthy that the results align with real-world expectations, providing a sensible justification for the price increases associated with the identified factors.

## Conclusion

---

The King County grading system used in this dataset reflects the quality of the houses, ranging from grade 1 to 13. The model results indicate that Grades 11, 10, and 9 are the top three in the ranking system, aligning with the realistic expectations of higher-quality renovations. Renovating a house to attain Grade 11, 10, or 9 not only increases its value but also suggests that additional improvements can further enhance the property's worth.

(source: <https://info.kingcounty.gov/assessor/esales/Glossary.aspx?type=r>  
(<https://info.kingcounty.gov/assessor/esales/Glossary.aspx?type=r>))

In summary, the primary drivers for maximizing a house's market value include:

1. Investing in high-quality materials and luxury options during renovation.
2. Securing a house with a waterfront view, if feasible.
3. Adding more bathrooms.